

# Advanced Topics on Artificial Intelligence

Alban Grastien

The Australian National University

Second Semester, 2020

# How to Represent MDPs

- MDPs can include a huge number of states.
- It is impossible to enumerate all these states, all the transitions, without making a mistake, even for smaller problems.
- Yet, the problem can often be simply described.
- Think game of go:
  - The rules can be written on a small page.
  - The number of states is astronomical.

# What is a Good Representation

- One goal is to come up with general-purpose planners
  - Cf. Atari games
- A good representation should be easy to write
  - Compact
  - Intuitive (in particular for non-experts)
  - No pit falls!
- A good representation should make the planning task easy.  
That includes:
  - Expose constants and mutexes
  - Make it easy to compute heuristics

# Probabilistic Planning Domain Description Language PPDDL

H. Younes and M. Littman. *PPDDL: The probabilistic planning domain definition language*. Technical Report. 2004.

# Probabilistic Planning Domain Description Language

```
(define (domain test-domain)
  (:requirements :typing :equality
    :conditional-effects :fluents)
  (:types car box)
  (:constants goldie - car)
  (:predicates (parked ?x - car) (holding ?x - box)
    (in ?x - box ?y - car))
  (:functions (fuel-level ?x - car))
  (:action load :parameters (?x - box ?y - car)
    :precondition (and (holding ?x) (parked ?y))
    :effect (and (in ?x ?y)(forall (?z - car)
      (when (not (= ?z ?y))(not (in ?x ?z))))))
  (:action refuel :parameters (?x - car)
    :precondition (< (fuel-level ?x) 10)
    :effect (increase (fuel-level ?x) 1))
)
```

# Types and Objects

- *Types* form a hierarchy of objects, the most abstract of which is `object`.
  - for instance `truck` and `plane` are subtypes of `vehicle`.
  - `location` is also a type.
  - both `vehicle` and `location` are subtypes of `object`.
- *Objects* refer to real-world instances of the types.
  - `truck_1` and `truck_2` are instances of `truck`, and therefore of `vehicle`.
  - `airport_1` and `warehouse_1` are instances of `location`.

# Predicates and Facts

- *Predicates* are properties that link objects of given types. They have a signature that indicates what type of objects they link:
  - `at_loc( ?t - vehicle ?l - location)` is a predicate with signature `vehicle,location`.
- *Facts* are instantiation of predicates with objects:
  - `at_loc(truck_1, warehouse_1)` models the fact that the truck 1 is in the first warehouse.

# Action Schemas and Actions

- *Action Schemas* are general definitions of actions that can be instantiated.
- An action schemas includes a name, a signature, a precondition, and a list of effects (more about effects later).

```
(:action drive
  :parameters (?t - truck ?lstart ?lend - location)
  :precondition (and (at_loc ?t ?lstart) (road ?lstart ?lend))
  :effect (and (not (at_loc ?t ?lstart)) (at_loc ?t ?lend))
)
```

- Action `drive(truck_1,warehouse_1,airport_1)` is an instantiation of the `drive` schema.



# Functions and Rewards

- *Functions* are used to represent numeric variables.

Warning: numeric variables make state space infinite, and problems undecidable in general.

- Functions are defined by schemas:

```
(:functions (fuel-level ?t - truck))
```

- Functions can be called in conditions: `(> (fuel-level ?t) 100)`

- Functions can be updated via action effects `assign`, `scale-up`,

`scale-down`, `increase`, `decrease`:

```
(increase (fuel-level ?t) 100)
```

- The *reward* function is accessed by the function `reward`.

# Conditional Effects

- Action preconditions and *conditional effects* are different:
  - An action precondition indicates when an action is applicable.
  - A conditional effect indicates when some effect applies.

The distinction is important when the current state is only partially known.

Example:

- Can you perform “move forward” if there is a wall in front of you?
- A conditional effect is represented by the keyword **when**:  

```
(when (<= (fuel-level ?t) 0) (stuck ?t))
```

# Uncertainty

- *Probabilistic effects* are represented as

`(probabilistic  $p_1$   $e_1$  ...  $p_k$   $e_k$ )`

where each  $p_i$  is a probability so that  $\sum_i p_i \leq 1$ .

- When the effect is supposed to trigger, then effect  $e_i$  triggers (and  $e_j$  does not if  $j \neq i$ ) with probability  $p_i$ .
- With probability  $1 - \sum_i p_i$ , no effect triggers.

# Initial State and Goal

- *Initial state* are defined as a conjunction of probabilistic facts.

```
(:init (and (at_loc truck_1 warehouse_1)
            (probabilistic .5 (at_loc plane_1 airport_1)
                          .5 (at_loc plane_1 airport_2))
            ...
))
```

- The *goal* is a formula on the facts and fluent values.


```
(:goal (or (> (fuel-level truck_1) 20)
           (at_loc plane_1 warehouse_1))
)
```

# Relational Dynamic Influence Diagram Language RDDL

S. Sanner, *Relational dynamic influence diagram language (RDDL): language description*, Technical Report, 2010.

Slides heavily inspired by Scott Sanner's Technical Report

# Influence Diagram



Left:

- Circles: state fluents
- Square: action fluents

Middle and right:

- Circles: intermediate & next state fluents
- Diamond: reward

# Define Variables

```
domain prop_dbn {
  requirements = { reward-deterministic };

  pvariables {
    // State fluents
    p : { state-fluent , bool , default = false };
    q : { state-fluent , bool , default = false };
    r : { state-fluent , bool , default = false };
    // Action fluents
    a : { action-fluent , bool , default = false };
    // Intermediate fluents
    i1 : { interm - fluent , int , level = 1 };
    i2 : { interm - fluent , int , level = 2 };
    // Observable fluents
    o1 : { observ - fluent , bool };
    o2 : { observ - fluent , real };
  };
};
```

# Define Values

```
...
cpfs {
  p' = if ( p ^ r ) then Bernoulli (.9) else Bernoulli (.3);
  ...

  i1 = KronDelta(p + r + q)

  ...

  o1 = Bernoulli ( ( p + q + r )/3.0 );
}

reward = p - r + i2
}
```



# Define Problem Instance

```
instance inst_dbn {  
    domain = prop_dbn2 ;  
    init-state { p ; r ; };  
    max-nondef-actions = 1;  
    horizon = 20;  
    discount = 0.9;  
}
```

# Lifted (Parametrised) Representation (1/2)

```
domain game_of_life {  
  ...  
  types {  
    x_pos : object ;  
    y_pos : object ;  
  };  
  
  pvariables {  
    alive ( x_pos , y_pos ) :  
      { state-fluent , bool , default = false };  
  };  
  
  cpfs {  
    alive'(?x ,?y ) = ...  
  }  
}
```

## Lifted (Parametrised) Representation (2/2)

```
non-fluents game2x2 {
  domain = game_of_life ;
  objects {
    x_pos : { x1 , x2 };
    y_pos : { y1 , y2 }; };
  non-fluents {
    // can define constants here
  };
}

instance is1 {
  domain = game_of_life ;
  non-fluents = game2x2 ;
  init-state {
    alive ( x1 , y1 );
    alive ( x2 , y2 ); };
  ...
}
```

# Scikit-decide

<https://airbus.github.io/scikit-decide/>

# Description

- Scikit-decide is a framework for decision making under uncertainty
- Compatible with AI gym

# How to Define an MDP (1)

<https://airbus.github.io/scikit-decide/guide/codegen.html>

```
class D(MDPDomain):
    T_state = State
    T_observation = T_state
    T_event = Action
    T_value = float
    T_info = None

class MyDomain(D):
    def _get_transition_value(self, memory: D.T_state,
        action: D.T_event, next_state: Optional[D.T_state] = None)
        -> TransitionValue[D.T_value]: pass

    def _get_next_state_distribution(self, memory: D.T_state,
        action: D.T_event)
        -> DiscreteDistribution[D.T_state]: pass
```

# How to Define an MDP (2/2)

<https://airbus.github.io/scikit-decide/guide/codegen.html>

```
def _is_terminal(self, state: D.T_state)
    -> bool: pass

def _get_applicable_actions_from(self, memory: D.T_state)
    -> Space[D.T_event]: pass

def _get_action_space_(self)
    -> Space[D.T_event]: pass

def _get_initial_state_(self)
    -> D.T_state: pass

def _get_observation_space_(self)
    -> Space[D.T_observation]: pass
```