

Advanced Topics on Artificial Intelligence

Alban Grastien

The Australian National University

Second Semester, 2020

Moving on

- Yesterday:
 - Policy Iteration (PI)
 - how to create/represent MDPs
- Today: more sophisticated algorithms
 - RTDP
 - LAO*

Real-Time Dynamic Programming RTDP

and variants

Asynchronous VI (aka Find & Revise)

ASYNCHRONOUS VI

- Choose an arbitrary value function $V : S \rightarrow \mathbb{R}$
- **repeat**
 - Choose state s
 - $V[s] := \min_{a \in A(s)} \sum_{s' \in S} P(s, a, s') \times \left(C(s, a, s') + \gamma V[s'] \right)$
- **until** some condition holds

Compared to Value Iteration:




- At each iteration, VI backs up all states ($V^{t+1} := BV$)
- At each iteration, A-VI backs up the value for only one state


Correctness

When is Asynchronous VI correct?

- 1 If no state is starved, then Asynchronous VI converges to V^*
 - A state is **starved** if it will eventually never be backed up
- 2 The states that are not reachable do not need to be backed up
- 3 If the value of state s does not change when backed up, you may also not back it up
- 4 If a state is provably not reachable under the optimal policy, then it does not need to be backed-up (more on that later).


Little Robot: Asynchronous VI

4	0	0	0	0	0
3	0	0	1	0	0
2		0		0	 100
1	0	50	0	0	0
0	0	0	0	0	0
	A	B	C	D	E

- Reward of getting to  = 100
- Discount factor = .9
- The value of which state should we rather backup?


Little Robot: Asynchronous VI

4	0	0	0	0	0
3	0	0	1	0	0
2	0	0		0	100
1	0	50	0	0	0
0	0	0	0	0	0
	A	B	C	D	E

- Reward of getting to  = 100
- Discount factor = .9
- The value of which state should we rather backup?


Topological VI

- Identifies connected components.
- Backs up the value of the inner components first.




Topological VI

- Identifies connected components.
- Backs up the value of the inner components first.



Topological VI

- Identifies connected components.
- Backs up the value of the inner components first.



- Occurs in problems with limited resources \rightarrow cf. Tinsley vs Chinook

English Draughts, aka American Checkers



(source Wikipedia)

- A game where you try to take all pieces from the opponent.
- If a piece is taken, it cannot reappear.



→ The positions with k pieces forms one or several connected components.


- The game was solved in 2007 at the Uni. of Alberta. It took 18 years.
- The researchers who solved it started from positions with 4 pieces, up.

Improving Asynchronous VI: Monotonicity

- $V(s) \leq V^*(s)$ for all $s \Rightarrow V(s) \leq BV(s) \leq V^*(s)$ for all s
- $V(s) \geq V^*(s)$ for all $s \Rightarrow V(s) \geq BV(s) \geq V^*(s)$ for all s
- The **policy graph** of π is the set of states reachable while following π .
- If
 - 1 the state s does not appear in the policy graph of π_V and
 - 2 $V(s') \leq V^*(s')$ for all states s' (in a minimisation context)then it is not necessary to backup this state (currently).


Little Robot: Monotonicity

4	4.62	3.75	3	2.5	2
3	5	3.5	2.5	1.5	1
2	 5.08	4.5		1.5	
1	5.93	4.87	3.75	2.5	1
0	6	5	4	3	2
	A	B	C	D	E

- Cost once in  = 0
- Cost of actions: 1


Little Robot: Monotonicity

4	4.62	3.75	3	2.5	2
3	5	3.5	2.5	1.5	1
2	5.68	4.5		1.5	0
1	5.93	4.87	3.75	2.5	1
0	6	5	4	3	2
	A	B	C	D	E

- Cost once in  = 0
- Cost of actions: 1
- Purple: policy graph

Little Robot: Monotonicity

4	4.62 4.5	3.75 3.65	3 2.75	2.5 2.25	2 1.5
3	5 4.8	3.5	2.5	1.5	1
2	5.68	4.5		1.5	0
1	5.93 5.5	4.87	3.75	2.5	1
0	6 5	5 4.5	4	3	2
	A	B	C	D	E

- Cost once in  = 0
- Cost of actions: 1
- Purple: policy graph
- Red: Estimates of the policy value for non-policy states. If we know that the red values are underestimate, it is not necessary to backup these states.

Real-Time Dynamic Programming

- RTDP builds on these properties.
- It only explores and backs up the current policy graph.
- Using an admissible heuristic, it can safely ignore the other states.

Real-Time Dynamic Programming

RTDP

- $V := h$
- Perform n trials

Trial:

- $s := s_0$
- **while** $s \notin G$
 - Backup state s :

$$V(s) := \min_{a \in A(s)} \sum_{s' \in S} P(s, a, s') \times \left(C(s, a, s') + \gamma V(s') \right)$$
 - Choose greedy action:




$$a := \arg \min_a \sum_{s' \in S} P(s, a, s') \times \left(C(s, a, s') + \gamma V(s') \right)$$
 - $s' := \text{simulate}(s, a)$ // i.e., outcome of a is randomised

Properties

- A heuristic is **admissible**
 - if it underestimates when you try to minimise cost
 - if it overestimates when you try to maximise reward.
- If the heuristic h is admissible, then RTDP converges towards the optimal policy

Illustrating RTDP




- Start with an **admissible heuristic**.

4	2	3	2	2	2
3	2	3	2	1	1
2		3		1	
1	2	3	2	1	1
0	2	3	2	2	2
	A	B	C	D	E


Illustrating RTDP

- Start with an **admissible heuristic**.

- Start trial 1:
 - Update A2, choose best action (here: Up or Down), and simulate it


4	2	3	2	2	2
3	2	3	2	1	1
2	 3	3		1	
1	2	3	2	1	1
0	2	3	2	2	2
	A	B	C	D	E

Illustrating RTDP




- Start with an **admissible heuristic**.
- Start trial 1:
 - Update A2, choose best action (here: Up or Down), and simulate it
 - Update A3, and choose best action (here: Up), and simulate it

Illustrating RTDP



- Start with an **admissible heuristic**.
- Start trial 1:
 - Update A2, choose best action (here: Up or Down), and simulate it
 - Update A3, and choose best action (here: Up), and simulate it
 - Update A4, and choose best action (here: Down or Right), and simulate it

Illustrating RTDP



- Start with an **admissible heuristic**.

- Start trial 1:

- Update A2, choose best action (here: Up or Down), and simulate it
- Update A3, and choose best action (here: Up), and simulate it
- Update A4, and choose best action (here: Down or Right), and simulate it
- Update A3, and choose best action (here: Down or Right), and simulate it (here, leads to B2)

RTDP's Performance



- RTDP is a “good” any-time algorithm:
 - Compared to VI or PI, it finds a good (sub-optimal) policy early
- By using the greedy policy, RTDP concentrates on the states that are likely to be in the optimal policy graph
- By using simulation, RTDP concentrates on the states that are more likely to actually be reached

L-RTDP¹

- Asynchronous VI is very efficient when it does not backup the states that have a good estimate.
- A state is **solved** if all states s' in the greedy space rooted at s have a Bellman error below ϵ .
- Backing up these states will not improve (significantly) the policy, while being time-consuming.
- L-RTDP identifies these solved states, and avoids them when simulating the effect of the action.
- Doing this is **not** unbiased (as opposed to Reinforcement Learning).


¹B. Bonet and H. Geffner. "Labeled RTDP: improving the convergence of real-time dynamic programming". In: *13th International Conference on Automated Planning and Scheduling*. 2003, pp. 12–21.

Example

4	4.5	3.65	2.75	2.25	1.5
3	4.8	3.48	2.49	1.49	1
2		4.48		1.5	
1	5	4.5	3.4	2.2	1
0	5	4.5	4	3	2
	A	B	C	D	E

- Current values


Example


4	4.5	3.65	2.75	2.25	1.5
3	4.8	3.48	2.49	1.49	1
2		4.48		1.5	0
1	5	4.5	3.4	2.2	1
0	5	4.5	4	3	2
	A	B	C	D	E

- Current values

-  Solved states

Example

4	4.5	3.65	2.75	2.25	1.5
3	4.8	3.48	2.49	1.49	1
2		4.48		1.5	0
1	5	4.5	3.4	2.2	1
0	5	4.5	4	3	2
	A	B	C	D	E

- Current values
-  Solved states
- Because B2 is solved, the effect of Right in A2 will always lead to B1 during simulation.

Remark

- Determining what states are solved is not a trivial task
- Search for Tarjan's algorithm

Bounding the policy

- It sometimes happens that the policy is not settled in parts of the greedy space that have a small probability of being reached.
- “What **will** I do if I end up in this situation that is unlikely to happen?”
 - I don't need to answer this question to decide the optimal action **now**.

Bounded-RTDP²

- Bounded-RTDP keeps two value functions V_ℓ and V_u such that $V_\ell(s) \leq V^*(s) \leq V_u(s)$ for all states.
 - Remember: the monotonicity property works if V is an upper bound or a lower bound
- Bounded-RTDP avoids states such that $(V_u(s) - V_\ell(s))/V_u(s) \leq \tau$ for some τ .
- Problem: How to compute (useful) upper bounds of $V^*(s)$?
 - It's harder than for admissible heuristics
 - Look at the paper for hints

²H. B. McMahan, M. Likhachev, and G. Gordon. "Bounded real-time dynamic programming: RTDP with monotone upper bounds and performance guarantees". In: *22nd International Conference on Machine Learning (ICML-05)*. 2005, pp. 569–576. 